

A recorded tool-use example: Coastal

Synthetic September 2026 office returns; use the highest approved revision.

1. Request the source

```
read_return  
office_id = Coastal
```

2. Observe the source rows

Rev.	Status	Issued	Ineligible	Complete
r1	approved	200	20	90
r2	approved	250	50	110
r3	<i>draft</i>	250	40	150

All counts are in cases. The note says r3 awaits approval.

3. Submit selected values

```
calculate_rates  
row: Coastal:r2  
multiplier: 1  
issued: 250; ineligible: 50  
completed: 110
```

6. Check the source

latest approved row: r2
counts and units agree
citations resolve

5. Return a finding

Coastal:r2
Coastal:definition
200 eligible; 110 completed
55%

4. Read the tool result

200 eligible cases
110 completed cases
55% completion

check all three offices and reducer totals

Validated pooled result

$$\begin{aligned}\text{completed} &= 700 + 320 + 110 = 1130, & \text{eligible} &= 1000 + 400 + 200 = 1600 \\ \text{pooled completion} &= 100 \times \frac{1130}{1600} = \mathbf{70.625\%}\end{aligned}$$

Recorded 6 October 2026 with GPT-6 Luna. Calls, observations, and outputs are abbreviated from the saved trace.
The calculator computes submitted counts; separate checks verify source selection and copied values.

Key

records / data

learned model / policy

LLM signal

labels / evaluation