

FEVER: finding evidence and testing claim support

Frozen slice: 12 claims (6 SUPPORTS, 6 REFUTES),
10 evidence sentences from 9 Wikipedia pages.

Separate retrieval illustration

12 claims
as search queries

BM25 ranks the 10 sentences
using page title and text

Annotated evidence
ranks first: 12/12

This small corpus is built from the selected annotations. It is not a Wikipedia-wide search test.

Recorded support task: annotated evidence supplied directly

Shared evidence input: page A_Milli, sentence 0

“A Milli”, abbreviated occasionally as “Milli”, is a song by American hip hop recording artist Lil Wayne.

Unstructured Wikipedia text; one supplied sentence per claim–evidence row.

Claim: FEVER 70373

A Milli is a song by a recording artist.

Claim: FEVER 143488

A Milli is a church.

LOTUS `sem_filter`: one decision per supplied claim–evidence row

The claim {claim} is supported by the supplied Wikipedia evidence {evidence}.

GPT-6 Luna returns a Boolean answer. The full instruction permits background knowledge to clarify terms; contradictory or insufficient evidence yields False.

70373: True

Retain the claim–evidence row.

143488: False

Exclude the claim–evidence row.

Evaluate against FEVER annotations, which are withheld from the model

70373: SUPPORTS; 143488: REFUTES. Both recorded decisions agree.

Across all 12 claims: 6 True and 6 False; 12/12 agree with the annotations.

The saved decisions preserve excluded rows for scoring.

Sources: FEVER development annotations (Thorne et al., 2018);
Wikipedia contributors, *A Milli*, June 2017 dump (CC BY-SA 3.0).

This selected slice contains no NOT ENOUGH INFO claims. False means the support predicate failed.